

noSQL Databases

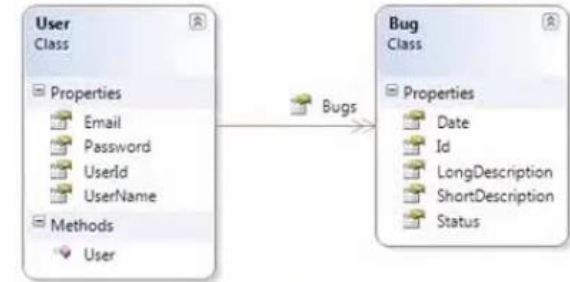
What is NoSQL?




Next Generation Databases

Not **O**nly SQL

Not Only SQL



Non – Relational



Distributed Architecture



Open Source



Horizontally Scalable

What is NoSQL?

Schema – Free !

Schema - Free



**Can Manage Huge
Amount of Data**



Easy – Replication



**Can be implement on
Commodity Hardware's**

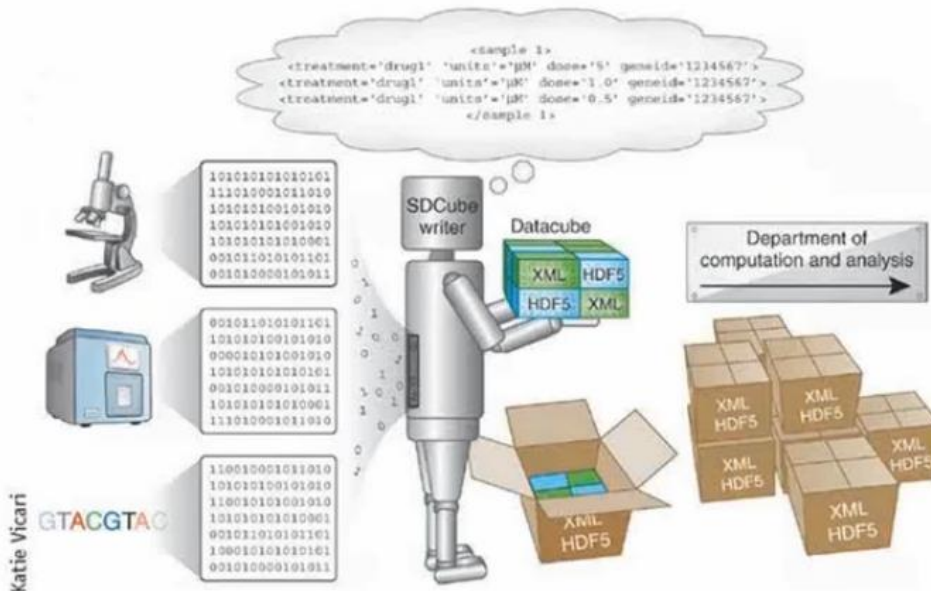


Simple API



**~ 150 No SQL Database
are there in Market**

Why NoSQL?

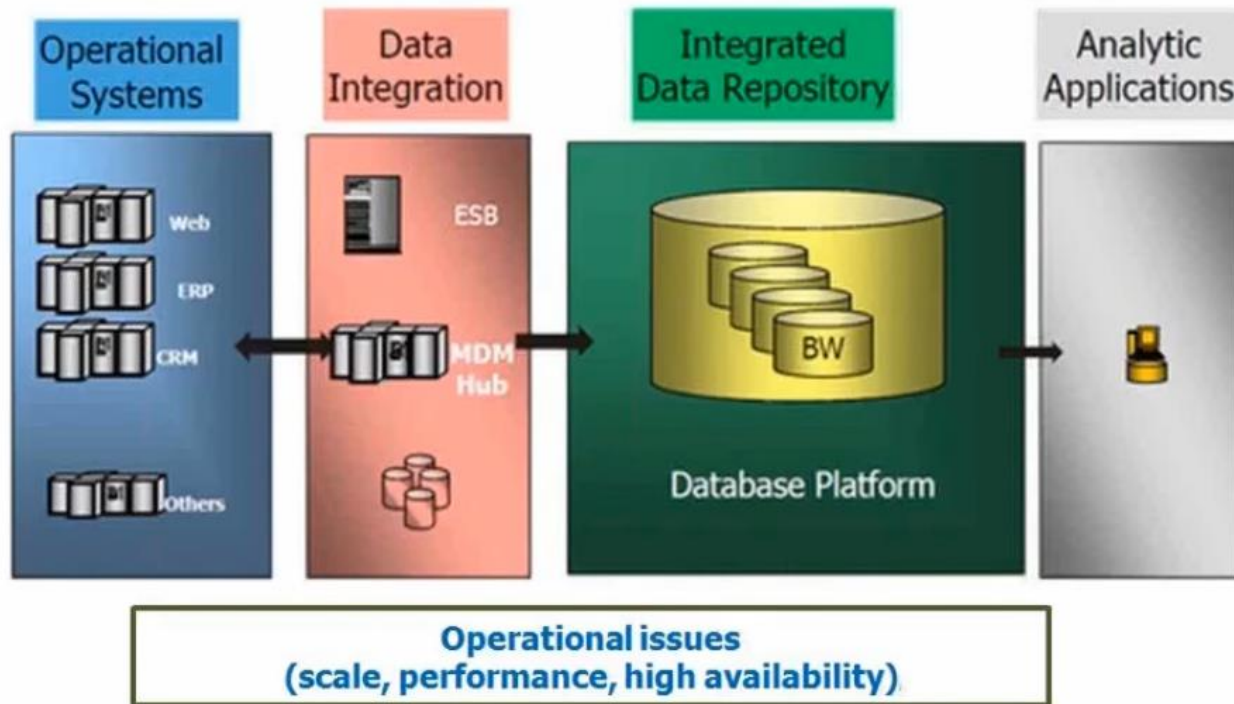


Nature of data

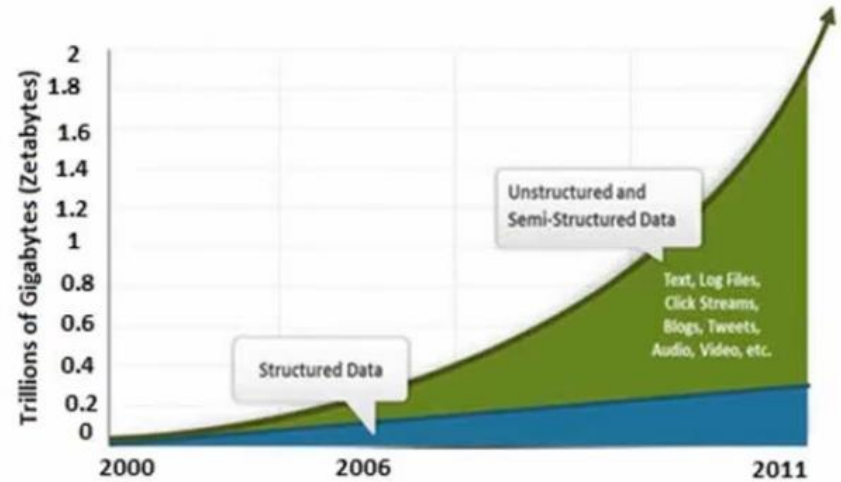



**Application development
(high coding velocity & agility)**

Why NoSQL?



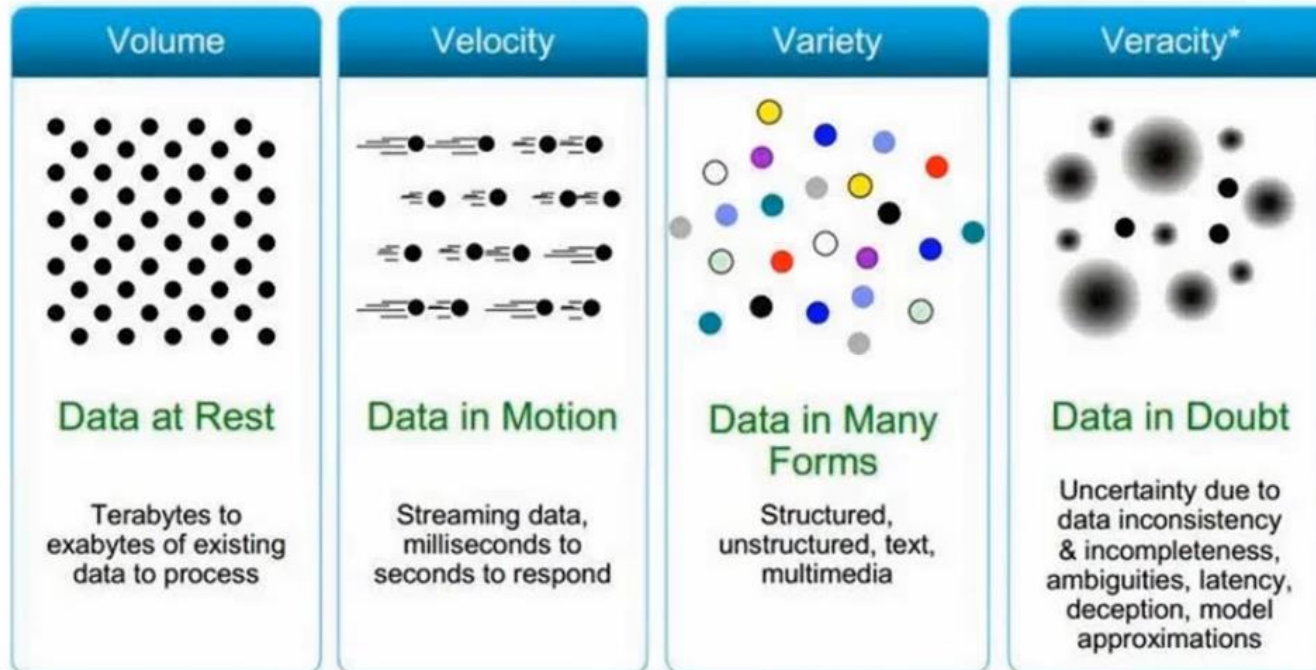
Why NoSQL?



Data warehousing and analytics

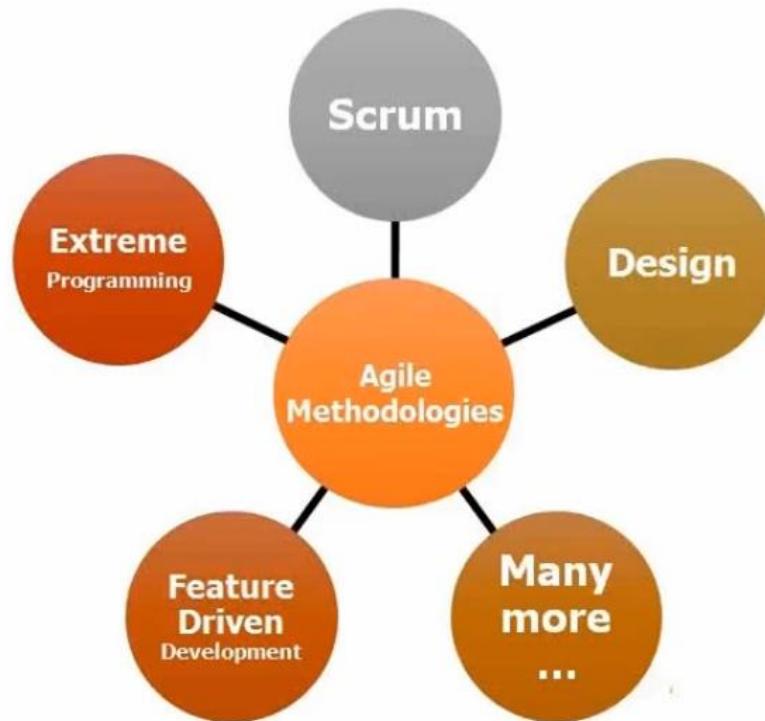
Benefits of NoSQL

- ✓ Large volumes of structured, semi-structured, and unstructured data



Benefits of NoSQL

- ✓ Agile development, quick changes, and frequent code pushes



Benefits of NoSQL

- ✓ Object-oriented programming that is easy to use and flexible



- ✓ Horizontal scaling instead of expensive hardware's



Categories of NoSQL Database

Document Base

- ✓ Document databases pair each key with a complex data structure known as a document.
- ✓ Documents can contain many different key-value pairs, or key-array pairs, or even nested documents.

Graph Store

- ✓ Graph stores are used to store information about networks, such as social connections.
- ✓ Graph stores include Neo4J and HyperGraphDB.

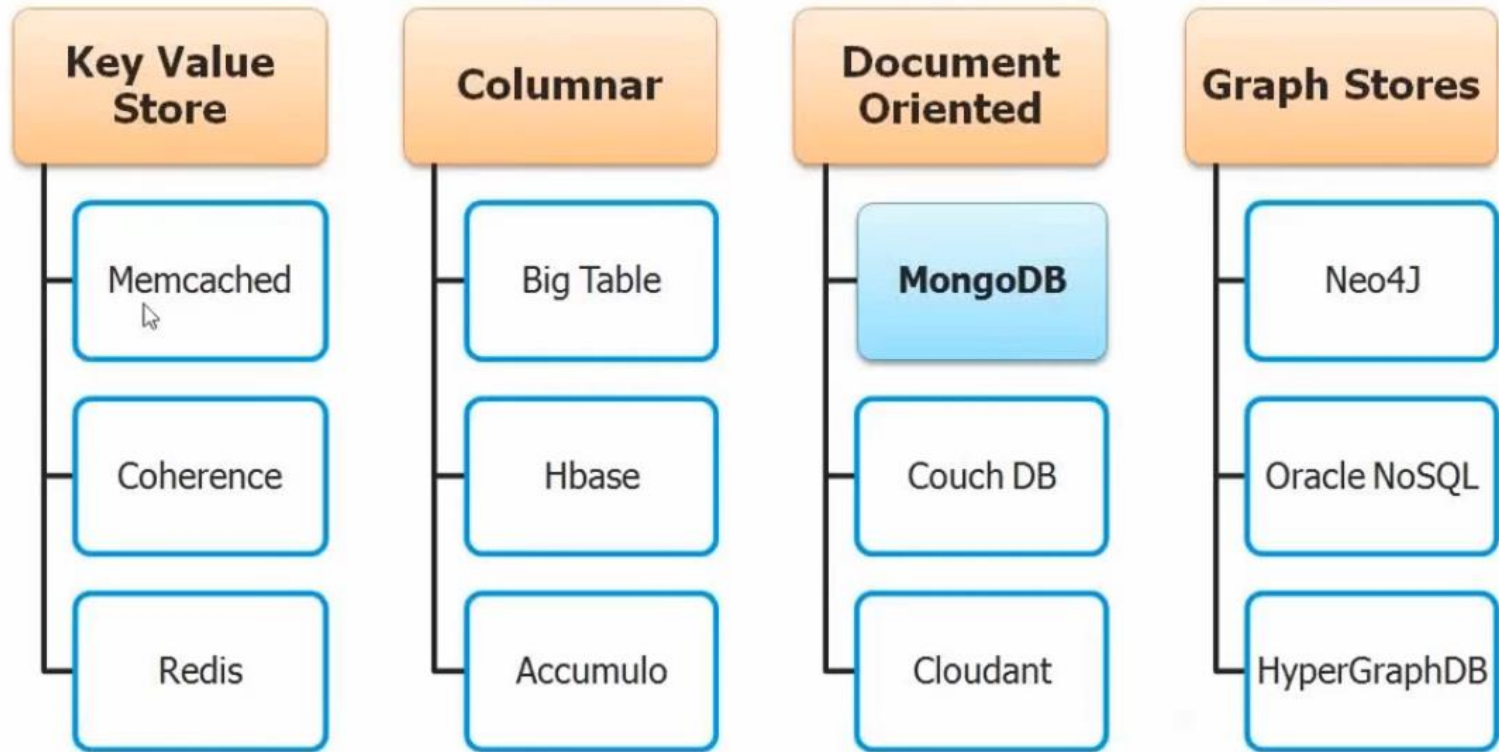
Key – value Stores

- ✓ Key-value stores are the simplest NoSQL databases.
- ✓ Every single item in the database is stored as an attribute name (or "key"), together with its value.

Wide Column Stores%

- ✓ Wide-column stores such as Cassandra and HBase are optimized for queries over large datasets, and store columns of data together, instead of rows.

Types of NoSQL Database



WHEN TO USE

NOSQL

AND WHEN TO USE

SQL

SQL

ATOMICITY

CONSISTENCY

ISOLATION

DURABILITY

noSQL

BASICALLY

AVAILABLE

SOFT STATE

EVENTUALLY

CONSISTENT

SQL

THE GOOD

High performance for transactions. Follow's ACID

Highly structured, very portable

Small amounts of data

*SMALL IS LESS THAN 500GB

Supports many tables with different types of data

Can fetch ordered data

Compatible with lots of tools

SQL

THE BAD

Complex queries take a long time

The relational model takes a long time to learn

Not really scalable

Not suited for rapid development

noSQL

THE GOOD

Fits well for volatile data

High read and write throughput

In general it's faster than SQL

Scales really well (Horizontally)

Rapid development is possible

noSQL

THE BAD

Key/Value pairs need to be packed/unpacked all the time

Still working on getting security for these working as well as SQL

Lack of relations from one key to another

tl;dr

SQL

works great, can't scale for large data

noSQL

works great, doesn't fit all situations

so use both, but think about when you want to use them!

NOSQL **USE CASES**

LARGE DATA VOLUMES

MASSIVELY DISTRIBUTED ARCHITECTURE

REQUIRED TO STORE THE DATA

GOOGLE, AMAZON, FACEBOOK, 100K SERVERS

EXTREME QUERY WORKLOAD

IMPOSSIBLE TO EFFICIENTLY DO JOINS AT THAT

SCALE WITH AN RDBMS

SCHEMA EVOLUTION

SCHEMA FLEXIBILITY IS NOT TRIVIAL AT A LARGE SCALE

BUT IT CAN BE WITH NO SQL